

# Kaushalendra Singh

+91 6398519072 | yadavkausha4a5@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

## EXPERIENCE

### EpicWise

AI Founding Engineer Intern

Remote (Canada Startup)

May 2026 – Present

- **Helped the founding team ship EpicWise's core context engine faster** by designing multi-step retrieval and reasoning pipelines using LangGraph.
- **Enabled secure third-party integrations with EpicWise** by building its public MCP server, including authentication, authorization, an agent harness testing agent, and a requirement-analysis agent.
- **Improved answer accuracy for complex user queries** by implementing a hybrid retrieval pipeline combining vector search (Pinecone) and graph-based retrieval (Neo4j), and authoring the system prompts, execution routers, and tool-calling schemas behind it.
- **Reduced production LLM errors** by building an evaluation pipeline that tracked tool correctness, answer relevancy, correctness, and hallucination metrics across agent runs using DeepEval, catching regressions before deployment.

### GRS Worker

Full Stack Developer Intern

Remote

Sep 2025 – Nov 2025

- **Supported reliable performance under concurrent load** by building a scalable backend using Node.js, Prisma, and PostgreSQL with Redis caching and secure authentication.
- **Reduced data errors by 40%** by implementing Zod validation and structured service-layer workflows.
- **Streamlined service deployment and customer communication** by integrating Twilio and QR-based systems, then containerizing and deploying all services with Docker and CI/CD pipelines.
- **Cut P95 API latency from ~800ms to ~400ms (50% reduction)** by optimizing database queries and the Redis caching strategy.

## OPEN SOURCE CONTRIBUTIONS

- **Improved stability for users of a widely-used desktop and mobile application (54K+ GitHub stars)** by fixing core issues across its desktop, mobile, and sync engine (React, Electron, Node.js), delivering merged and approved pull requests.
- **Built credibility on a large real-world codebase** by collaborating directly with project maintainers and following structured PR review workflows.

## PROJECTS

### Buildify Labs — Agentic AI Business Intelligence

[Live](#) | [GitHub](#)

React, Vite, FastAPI, Groq, HuggingFace, PostgreSQL, pgvector, Zustand

- **Turned plain-English questions into SQL-driven charts and explanations** by building an AI BI copilot on a judge-narrate-guarantee pipeline with LLM tool routing, sufficiency judging, and deterministic visual synthesis with grounding checks.
- **Grounded answers in each user's own documents** by implementing RAG over uploaded PDFs: text extraction and chunking, HuggingFace embeddings, and per-user-scoped pgvector cosine retrieval, returned as validated JSON and streamed live over SSE.
- **Made multi-turn conversations resolve without re-asking** by adding thread-scoped prior-turn context and persisted chat history, with context-window budgeting through row truncation, evidence ranking, and long-source summarization.

### Evident — Open-Source Extraction Layer for AI Agents

[Live](#) | [GitHub](#)

Python, MCP, REST API, Crawl4AI, Playwright, pytest, GitHub Actions

- **Gave AI agents a way to judge how much to trust extracted data** by building an open-source extraction and fetch layer that attaches a confidence score and method to every result, callable via MCP, REST, or a native Python SDK.
- **Made reliable extraction possible on sites without a hand-written parser** by implementing a tiered resilience ladder, static fetch escalating to JS-rendered fetch via Crawl4AI/Playwright, plus a versioned, community-contributable recipe system covering sites like Greenhouse, Lever, and Ashby.
- **Protected the fetch layer from server-side request forgery** by building an SSRF guard that resolves hostnames to IPs and blocks private, loopback, and cloud-metadata address ranges before any connection is opened.

### RepoGraph AI — VS Code Extension (50+ Downloads)

[Live](#) | [GitHub](#)

TypeScript, VS Code API, Groq, Ollama, Gemini, Anthropic, OpenAI

- **Helped developers understand unfamiliar codebases faster** by building a VS Code extension that scans workspaces and maps every file relationship into an interactive, force-directed dependency graph.
- **Gave users flexibility to choose their own AI provider** by integrating 5 AI providers (Groq, Ollama, Gemini, Anthropic, OpenAI) with custom model support and zero telemetry.
- **Made codebase Q&A more reliable** by implementing a multi-turn agent with full graph, summary, and file context, plus smart filtering across 20+ languages.

### MyAttendance

[Live](#) | [GitHub](#)

Next.js, MongoDB, Redis, PWA

- **Served 1300+ users reliably** by building a PWA with role-based access control and MongoDB queries optimized for high concurrency.
- **Reduced database load and improved uptime visibility** by implementing a Redis caching layer and structured error logging for production monitoring.

## AI ENGINEERING

- **Agentic Systems:** LangChain, LangGraph; ReAct & Plan-Execute patterns; tool calling; multi-step reasoning workflows.
- **RAG & Retrieval:** Embedding models, Pinecone/FAISS vector DBs, semantic search, context & memory management.
- **LLM Systems:** OpenAI/Gemini/Anthropic APIs with structured outputs, prompt engineering, state handling in AI workflows.
- **Model Development:** PyTorch for model experimentation and fine-tuning; Python for data and pipeline scripting.
- **Evaluation:** DeepEval for LLM output evaluation, regression testing, and quality benchmarking of AI agent responses.
- **MCP:** Designed and built MCP servers and integrations, auth, authorization, agent harness testing, and requirement-analysis agents.

## EDUCATION

### Harcourt Butler Technical University, Kanpur

Bachelor of Technology

CGPA: 7.0 / 10

## TECHNICAL SKILLS

- **Languages:** JavaScript/TypeScript, Python, SQL, C++
- **Frontend:** Next.js, React, Tailwind, Zustand
- **Backend:** Node.js, Express, FastAPI, Prisma, REST/GraphQL, Kafka, Redis
- **AI/ML:** LangChain, LangGraph, PyTorch, OpenAI, Gemini, Anthropic, MCP, Neo4j, FAISS, Pinecone, DeepEval
- **Database:** MongoDB, PostgreSQL, MySQL
- **DevOps:** Docker, Kubernetes, AWS, CI/CD, Logging & Monitoring